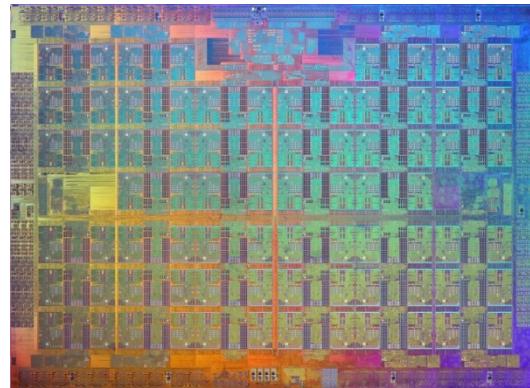
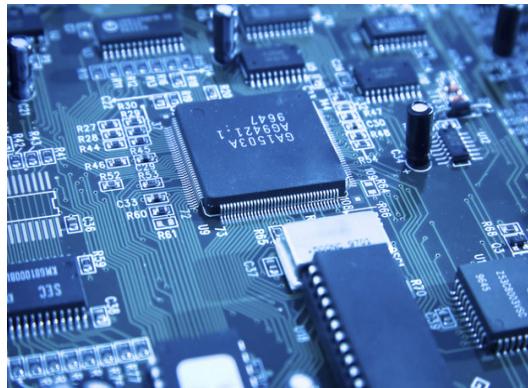


Welcome to 6.004!

Computation Structures



Spring 2020

6.004 Course Staff

Instructors



Silvina Hanono Wachman
silvina@mit.edu



Jason Miller
jasonm@csail.mit.edu

Teaching Assistants

Shahul
Alam



Alan
Cheng



Kendall
Garner



Domenic
Nutile



Brian
Chen



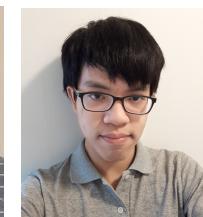
Miles
Dai



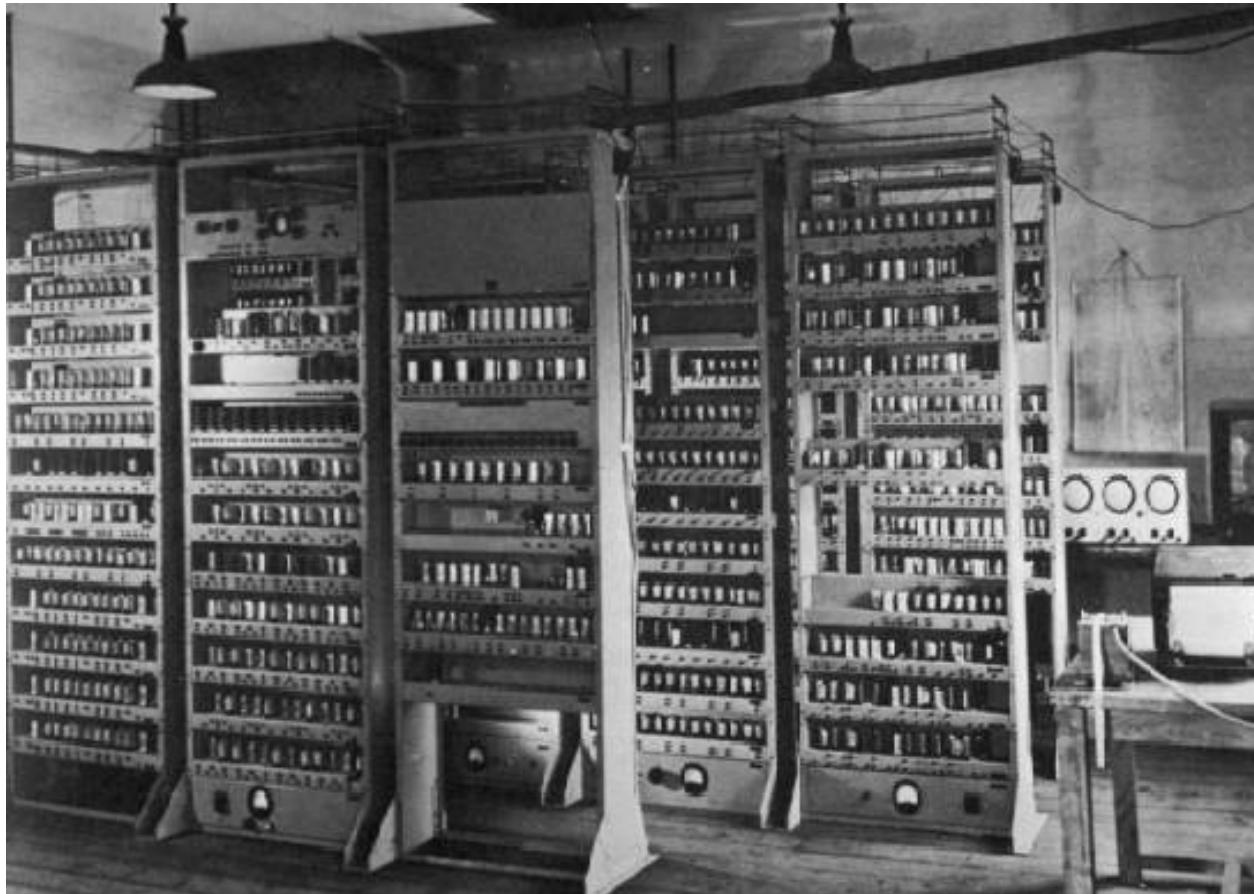
Min Thet
Khine



Boom
Tiankanon



Computing Devices Then...



ENIAC, 1943 30 tons, 200KW, ~1000 ops/sec
一打开房间的灯都会变暗

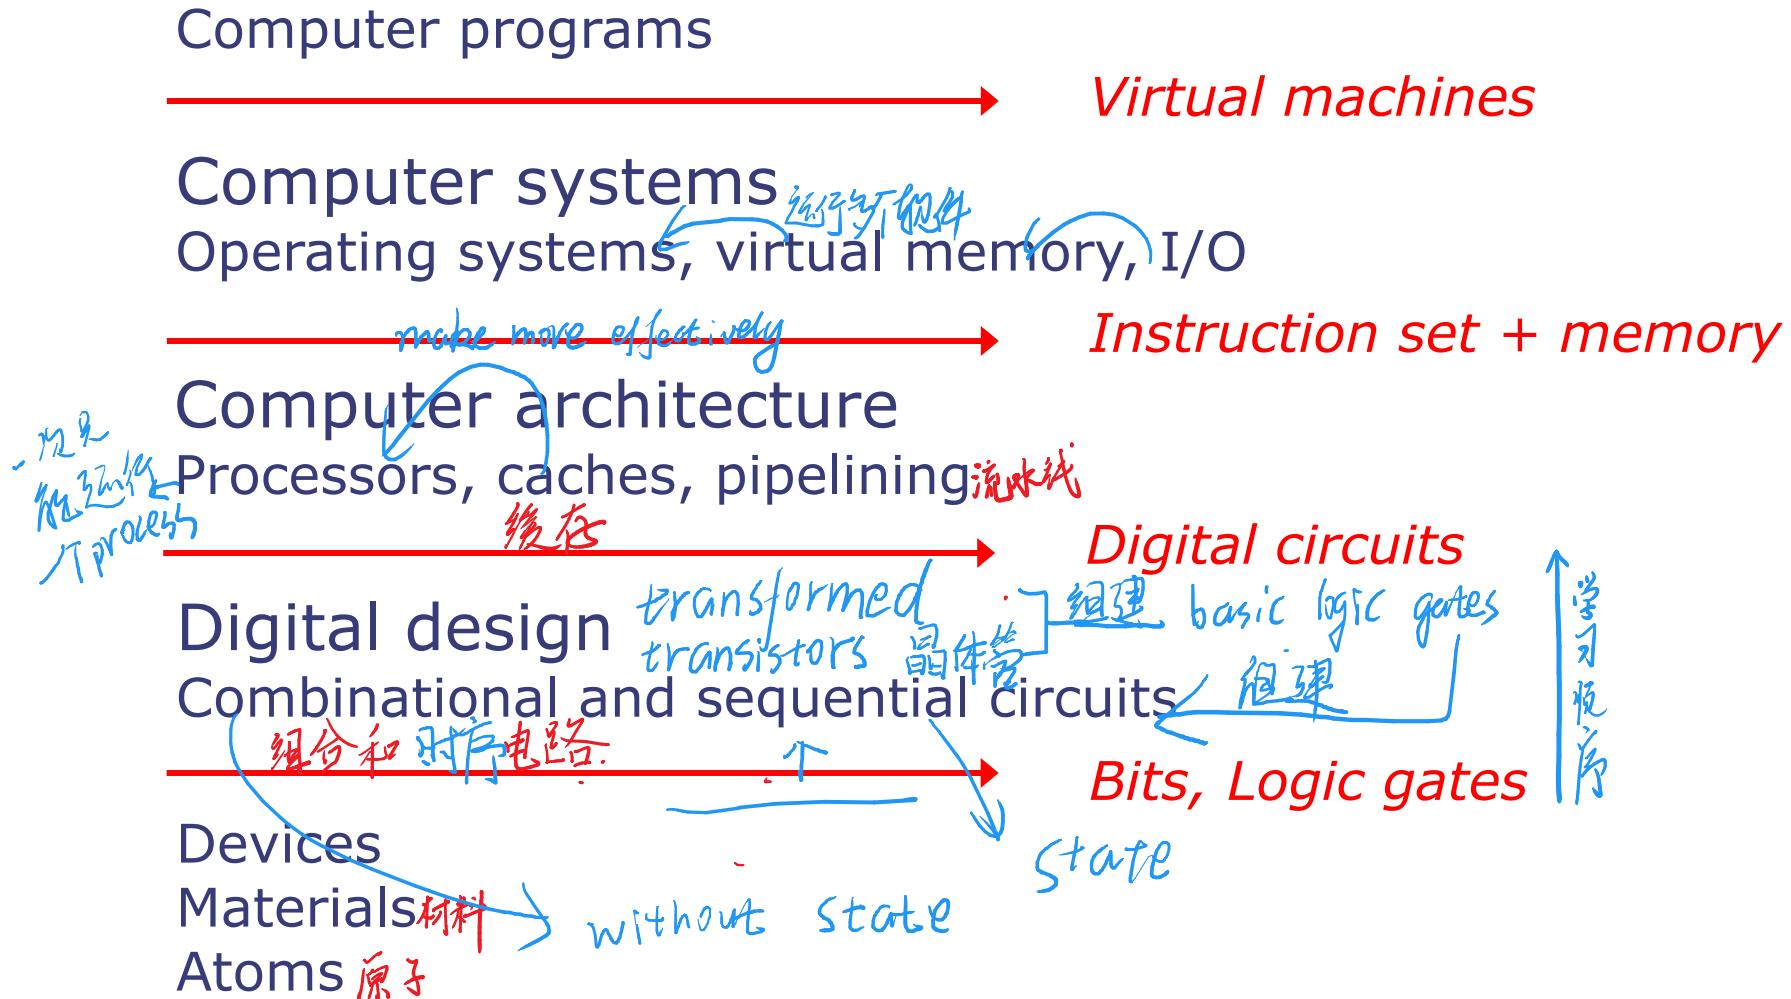
Computing Devices Now



Typical 2020 laptop
1kg, 10W, 10 billion ops/s

都使用到了 general-purpose computer

An Introduction to the Digital World



The Power of Engineering Abstractions

抽象
Abstraction

Good abstractions let us reason about behavior while shielding us from the details of the implementation.

推论
Corollary: implementation technologies can evolve while preserving the engineering investment at higher levels.

抽象
Abstraction

Leads to hierarchical design:

- Limited complexity at each level \Rightarrow shorten design time, easier to verify
- Reusable building blocks

Virtual machines

Instruction set + memory

Digital circuits

Bits, Logic gates

Course Outline

目录

- Module 1: Assembly language
 - From high-level programming languages to the language of the computer
- Module 2: Digital design
 - Combinational and sequential circuits
- Module 3: Computer architecture
 - Simple and pipelined processors
 - Caches and the memory hierarchy
- Module 4: Computer systems
 - Operating system and virtual memory
 - Parallelism and synchronization

机器语言

数字设计

体系结构

系统

Our Focus: Programmable General-Purpose Processors

- Microprocessors are the basic building block of computer systems
 - Understanding them is crucial even if you do not plan to work as a hardware designer
- Microprocessors are the most sophisticated digital systems that exist today
 - Understanding them will help you design all kinds of hardware

By the end of the term you will have designed a simple processor from scratch!

We Rely on Modern Design Tools

- We will use RISC-V, a simple and modern instruction set

架构

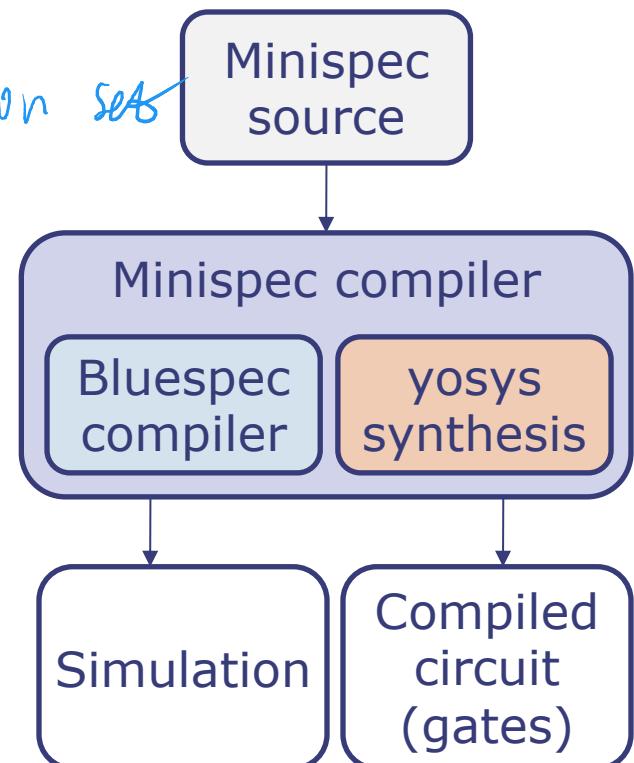
指令集

x86 Arm 等等 instruction sets

- We will design hardware using Minispec, a new hardware description language built for 6.004
 - Based on Bluespec, → 硬件描述语言
 - but heavily simplified

Minispec Minispec Minispec 6.004

硬件描述语言-Bluespec 硬件描述语言



Course Website – 6004.mit.edu

- Home: Announcements
- Information: Course information
- Material
- Labs
- Help
- User: View grades

.

Course Mechanics

- 2 lectures/week: slides, videos, and reference materials on website
- 2 recitations/week: work through tutorial problems using skills and concepts from previous lecture
- 7 mandatory lab exercises
 - Online submission + check-off meetings in lab
 - Due throughout the term (7 free late days meant to give you flexibility, cover short illnesses, etc.; see website)
- One open-ended design project
 - Due at the end of the term
- 3 quizzes: March 5, April 16, May 7 (7:30-9:30pm)
 - If you have a conflict, contact us to schedule a makeup

Recitation Mechanics

- 10 recitation sections on Wed & Fri
 - If you have a conflict with your assigned section, let us know your availability via piazza
- Recitation attendance is mandatory (worth 5% of your grade)
 - Everyone has 4 excused absences
- Recitations will review lecture material and problems associated with each lecture
 - We recommend you work on the problems before recitation
 - **Many quiz problems will be similar to recitation problems**

MIT 6.004 Recitation Mechanics
Quiz Problems

Grading

- 80 points from labs,
20 points from design project,
90 points from quizzes,
10 points from recitation attendance
- Fixed grade cutoffs:
 - A: Points ≥ 165
 - B: Points ≥ 145
 - C: Points ≥ 125
 - F: Points < 125 or not all labs complete

Online and Offline Resources

- The course website has up-to-date information, handouts, and references to supplemental reading:
<http://6004.mit.edu>
- We use Piazza extensively
 - Fastest way to get your questions answered
 - All course announcements are made on Piazza
- We will hold regular office hours in the lab (room 32-083) to help you with lab assignments, infrastructure, and any other questions

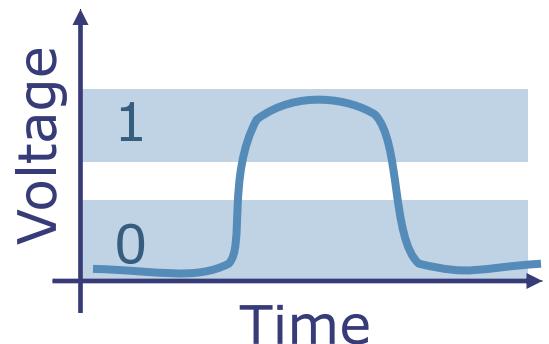
32-083 Combination Lock: _____

Binary Number Encoding and Arithmetic

二进制数的编码与运算

Digital Information Encoding

- Digital systems represent and process information using **discrete symbols** or digits
- These are typically binary digits (bits): 0 and 1
对应的逻辑电平的高低，元件的
(物理)
- We can implement operations like +, >, AND, etc. on binary numbers in hardware very efficiently
运算规则简单



Encoding Positive Integers

二进制数
Binary Number
比特位 Bit

It is straightforward to encode positive integers as a sequence of bits. Each bit is assigned a weight. Ordered from right to left, these weights are increasing powers of 2. The value of an N-bit number is given by the following formula:

$$v = \sum_{i=0}^{N-1} 2^i b_i$$

	2^{11}	2^{10}	2^9	2^8	2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0
	0	1	1	1	1	1	0	1	0	0	0	0

What value does 011111010000 encode?

$$\begin{aligned} V &= 0*2^{11} + 1*2^{10} + 1*2^9 + \dots \\ &= 1024 + 512 + 256 + 128 + 64 + 16 = 2000 \end{aligned}$$

Smallest number? 0

Largest number? $2^N - 1$

Hexadecimal Notation 对 binary 简写

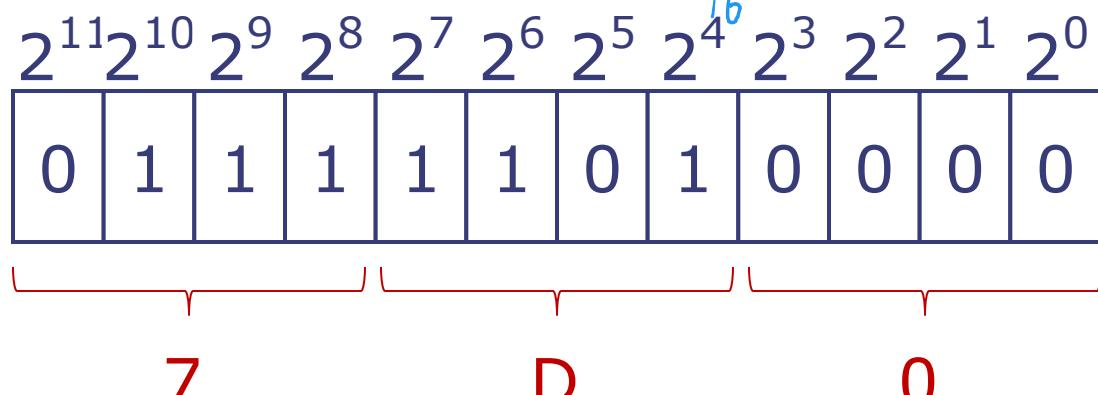
冗长的
易出错.
Long strings of bits are tedious and error-prone to transcribe, so we often use a higher-radix notation, choosing the radix so that it's simple to recover the original bit string.

A popular choice is to transcribe numbers in base-16, called hexadecimal. Each group of 4 adjacent bits is represented as a single hexadecimal digit.

相邻
 $(2^3 + 2^2 + 2^1 + 2^0) \underbrace{2}_{16}^4$

Hexadecimal - base 16

0000	-	0	1000	-	8
0001	-	1	1001	-	9
0010	-	2	1010	-	A
0011	-	3	1011	-	B
0100	-	4	1100	-	C
0101	-	5	1101	-	D
0110	-	6	1110	-	E
0111	-	7	1111	-	F



$$0b011111010000 = 0x7D0$$

Binary Addition and Subtraction

- Addition and subtraction in base 2 are performed just like in base 10

Base 10

$$\begin{array}{r} & \text{--- carry} \\ & 1 \\ 14 & \\ + & 7 \\ \hline 21 \end{array}$$

-1 — borrow

$$\begin{array}{r} \\ -1 \\ 14 \\ - 7 \\ \hline 07 \end{array}$$

Base 2

$$\begin{array}{r} 111 \\ 1110 \\ + 111 \\ \hline 10101 \end{array}$$

$$\begin{array}{r} -1-1-1 \\ 1110 \\ - 111 \\ \hline 0111 \end{array}$$

What does this mean?

$$-2^3 + 0b110$$

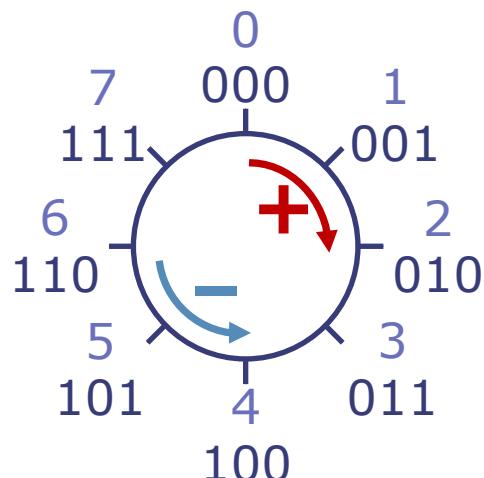
We need a way to represent negative numbers!

11101101
-1

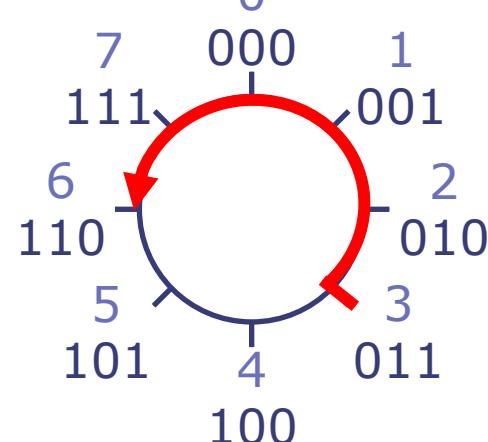
$$\begin{array}{r} 011 \\ - 101 \\ \hline ??? 110 \end{array}$$

Binary Modular Arithmetic

- If we use a fixed number of bits, addition and other operations may produce results outside the range that the output can represent (up to 1 extra bit for addition)
 - This is known as an **overflow** like to fixed number 3 bits
000 + 001 → (5+3)mod8
- Common approach: Ignore the extra bit
 - Gives rise to **modular arithmetic**: With N-bit numbers, equivalent to following all operations with $\text{mod } 2^N$
 - Visually, numbers “wrap around”:



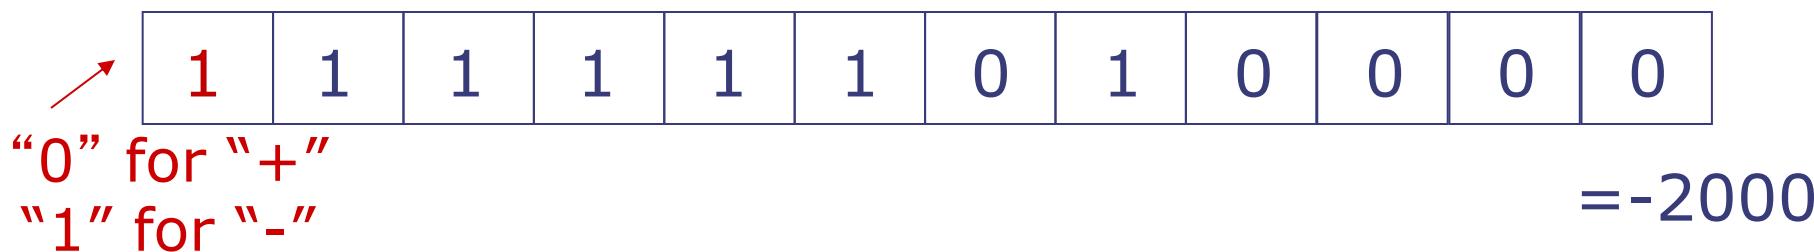
Example: $(3 - 5) \text{ mod } 2^3 ?$



Encoding Negative Integers

Attempt #1: Use a sign-magnitude representation for decimal numbers, encoding the sign of the number (using "+" and "-") separately from its magnitude (using decimal digits).

We could use the same approach for binary representations:



What issues might this encoding have?

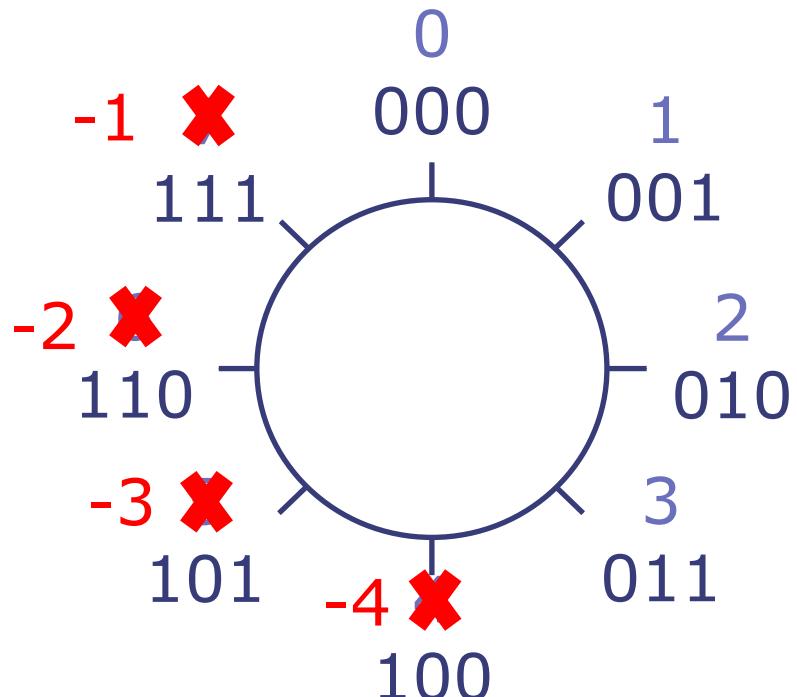
Two representations for 0 (+0, -0)

Addition and subtraction use different algorithms and are more complex than with unsigned numbers

Deriving a Better Encoding

Can you simply relabel some of the digits to represent negative numbers while retaining the nice properties of modular arithmetic?

Yes!



This is called two's complement encoding

Two's Complement Encoding

In two's complement encoding, the high-order bit of the N-bit representation has negative weight:

$$v = -2^{N-1}b_{N-1} + \sum_{i=0}^{N-2} 2^i b_i$$

最高位负，其余位正

正：与原码相同。
负：除符号位取反，加1。

- Negative numbers have “1” in the high-order bit
- *Most negative number?* $10\ldots0000 \quad -2^{N-1}$
- *Most positive number?* $01\ldots1111 \quad +2^{N-1} - 1$
- *If all bits are 1?* $11\ldots1111 \quad -1$

Two's Complement and Arithmetic

- To negate a number (i.e., compute $-A$ given A), we invert all the bits and add one

Why does this work?

$$\begin{aligned} -A + A &= 0 = -1 + 1 \\ -A &= (-1 - A) + 1 \\ &\quad \text{---} \\ &\quad \begin{matrix} 1 & \dots & 1 & 1 \end{matrix} \\ &\quad \frac{-A_{n-1} \dots A_1 A_0}{\overline{A}_{n-1} \dots \overline{A}_1 \overline{A}_0} \end{aligned}$$

Handwritten annotations: A blue 'Z' with a bar over it is on the left. A blue circle around '-1' is above the first equation. Blue arrows point from the circled '-1' to the circled '-1' in the second equation, and from the circled '-1' to the '1' in the third equation. A blue bracket under the '1 ... 1 1' row groups the first three digits. A blue '1' is written next to the fourth digit.

- To compute $A-B$, we can simply use addition and compute $A+(-B)$
 - Result: Same circuit can add and subtract!

Two's Complement Example

Compute $3 - 6$ using 4-bit 2's complement addition

- 3: 0011
- 6: 0110
- -6: 1010

$$\begin{array}{r} & \overset{1}{\cancel{0}} \\ & + \\ 0011 & + 1010 \\ \hline 1101 \end{array}$$

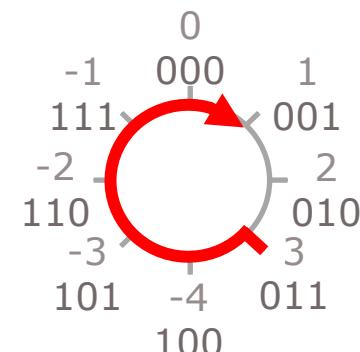
Compute $3 - 2$ using 3-bit 2's complement addition

- 3: 011
- 2: 010
- -2: 110

$$\begin{array}{r} 11 \\ 011 \\ + 110 \\ \hline 1001 \end{array}$$

What does this 1 mean?

Keep only last 3 bits



Zero crossing

Summary

- Digital systems encode information using binary for efficiency and reliability
- We can encode unsigned integers using strings of bits; long addition and subtraction are done as in decimal
- Two's complement allows encoding negative integers while preserving the simplicity of unsigned arithmetic

Thank you!

Next lecture:
Introduction to assembly and RISC-V